

# SEASONAL ARIMA APPLIED TO FASHION RETAIL

In this work, we have analysed data from fashion sales from 4 years in one category. The objective is sales forecasting for the following years. To achieve this goal, we have used Sarima Model as fashion sales is known to have an stationary performance. The steps followed were:

1. Import data already filtered and grouped by month
2. Visualize data
3. Analyze stationarity with Addfuller test
4. Analyze Autocorrelation and Partial Autocorrelation to see stationarity and seasonality
5. Diferentiation
6. Apply SARIMA model
7. Comparing results between the chosen parameters
8. Forecasting with new model

## 1. Import Data

```
In [12]: import numpy as np
import pandas as pd
import warnings
warnings.filterwarnings("ignore")
import matplotlib.pyplot as plt
%matplotlib inline
from statsmodels.tsa.stattools import adfuller
import statsmodels.api as sm
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
```

```
In [15]: df=pd.read_csv('sales-textiles3.csv')
```

```
In [16]: df.head()
```

Out[16]:

	Date	Sales
0	1997-01-01	16475
1	1997-02-01	26996
2	1997-03-01	41323
3	1997-04-01	31550
4	1997-05-01	8226

```
In [17]: ## Cleaning up the data  
df.columns=["Month", "Sales"]  
df.head()
```

Out[17]:

	Month	Sales
0	1997-01-01	16475
1	1997-02-01	26996
2	1997-03-01	41323
3	1997-04-01	31550
4	1997-05-01	8226

```
In [18]: # Convert Month into Datetime  
df['Month']=pd.to_datetime(df['Month'])
```

```
In [19]: df.head()
```

Out[19]:

	Month	Sales
0	1997-01-01	16475
1	1997-02-01	26996
2	1997-03-01	41323
3	1997-04-01	31550
4	1997-05-01	8226

```
In [20]: df.set_index('Month', inplace=True)
```

```
In [21]: df.head()
```


Out[21]:

	Sales
Month	
1997-01-01	16475
1997-02-01	26996
1997-03-01	41323
1997-04-01	31550
1997-05-01	8226

## 2. Visualize the Data

```
In [24]: df['Sales'].plot()
```

```
Out[24]: <matplotlib.axes._subplots.AxesSubplot at 0x224ae3ab320>
```




```
In [26]: # Here we divide data into train and test parts
# The test part will be the last year
y_train=df[:36]
y_test=df[36:]
```

```
In [ ]: # We can observe that date is Seasonal, as it responds to a fashion sector per
         formance.
```

## Auto Regressive Model

$$y_t = c + \phi_1 y_{t-1} + \phi_2 y_{t-2} + \cdots + \phi_p y_{t-p} + \varepsilon_t,$$

```
In [28]: from pandas.tools.plotting import autocorrelation_plot
autocorrelation_plot(y_train['Sales'])
plt.show()
```



```
In [ ]: #The example above shows positive first-order autocorrelation, where first order indicates that observations that are one apart are correlated, and positive means that the correlation between the observations is positive.
#As the points appear in a smooth snake-like curve, it means the correlation is positive
```

### 3. Analyze stationarity

```
In [27]: test_result=adfuller(y_train['Sales'])
#Ho: It is non stationary
#H1: It is stationary

def adfuller_test(sales):
    result=adfuller(sales)
    labels = ['ADF Test Statistic','p-value','#Lags Used','Number of Observations Used']
    for value,label in zip(result,labels):
        print(label+' : '+str(value) )
    if result[1] <= 0.05:
        print("strong evidence against the null hypothesis(Ho), reject the null hypothesis. Data has no unit root and is stationary")
    else:
        print("weak evidence against null hypothesis, time series has a unit root, indicating it is non-stationary ")


adfuller_test(y_train['Sales'])

ADF Test Statistic : -4.19604101801485
p-value : 0.000669141717223674
#Lags Used : 7
Number of Observations Used : 28
strong evidence against the null hypothesis(Ho), reject the null hypothesis.
Data has no unit root and is stationary
```

```
In [ ]: #The Adfuller test tell us that it is stationary to, so we don't need to differentiate
```

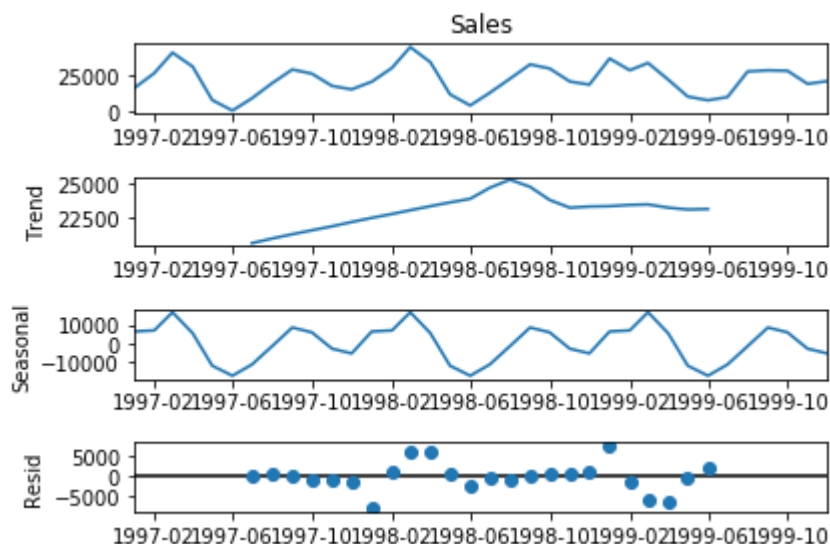
## 4. Analyze Autocorrelation and Partial Autocorrelation

```
In [29]: fig = plt.figure(figsize=(12,8))
ax1 = fig.add_subplot(211)
fig = sm.graphics.tsa.plot_acf(y_train['Sales'],lags=35,ax=ax1)
ax2 = fig.add_subplot(212)
fig = sm.graphics.tsa.plot_pacf(y_train['Sales'],lags=35,ax=ax2)
```



```
In [ ]: #With AFC we can visualize seasonality at first year. We can see the asesonl patron better in the following image
```


```
In [30]: ts_decomp=sm.tsa.seasonal_decompose(y_train['Sales'],model='additive')
ts_decomp.plot()
plt.show()
```



## 5. Differentiation

```
In [31]: y_train['Seasonal First Difference']=y_train['Sales']-y_train['Sales'].shift(15)
y_train['Seasonal First Difference'].plot()
#As we know that there is seasonality, we have first shift 12. But As we have seen there is no stationarity with this differentiation, we have then shift 3 more (15)
```

Out[31]: <matplotlib.axes.\_subplots.AxesSubplot at 0x224af775198>



```
In [ ]: #We have performed adfuller test until having stationarity
```

```
In [32]: test_result=adfuller(y_train['Seasonal First Difference'].iloc[15:])
def adfuller_test(sales):
    result=adfuller(sales)
    labels = ['ADF Test Statistic','p-value','#Lags Used','Number of Observations Used']
    for value,label in zip(result,labels):
        print(label+' : '+str(value) )
    if result[1] <= 0.05:
        print("strong evidence against the null hypothesis(Ho), reject the null hypothesis. Data has no unit root and is stationary")
    else:
        print("weak evidence against null hypothesis, time series has a unit root, indicating it is non-stationary ")

adfuller_test(y_train['Seasonal First Difference'].iloc[15:])
```

ADF Test Statistic : -6.351479265416526


p-value : 2.604405490424686e-08

#Lags Used : 1

Number of Observations Used : 19

strong evidence against the null hypothesis(Ho), reject the null hypothesis.  
Data has no unit root and is stationary

```
In [33]: fig = plt.figure(figsize=(12,8))
ax1 = fig.add_subplot(211)
fig = sm.graphics.tsa.plot_acf(y_train['Seasonal First Difference'].iloc[15:],
lags=20, ax=ax1)
ax2 = fig.add_subplot(212)
fig = sm.graphics.tsa.plot_pacf(y_train['Seasonal First Difference'].iloc[15:],lags=20,ax=ax2)
```



```
In [ ]: #We can not conclude something with the graphics,
#We will perform iterations so we will see best parameters for SARIMA
```

## 6. Applying SARIMA model

```
In [34]: y_train.index = pd.DatetimeIndex(y_train.index.values,
freq=y_train.index.inferred_freq)
```

```
In [ ]: #We have performed iterations with 1, 2 and 3 number of pdq. Here I leave the
iteration with number 1
```



In [36]: `import itertools`

```
p=d=q=range(0,2)
pdq=list(itertools.product(p,d,q))
seasonal_pdq=[(x[0],x[1],x[2],12) for x in list(itertools.product(p,d,q)) ]
print(seasonal_pdq)
print ('Examples of parameter combinations for Seasonal ARIMA...')
print('SARIMAX:{ }X{ }'.format(pdq[0],seasonal_pdq[0]))
```

```
[(0, 0, 0, 12), (0, 0, 1, 12), (0, 1, 0, 12), (0, 1, 1, 12), (1, 0, 0, 12),
(1, 0, 1, 12), (1, 1, 0, 12), (1, 1, 1, 12)]
Examples of parameter combinations for Seasonal ARIMA...
SARIMAX:(0, 0, 0)X(0, 0, 0, 12)
```

```
In [37]: metric_aic_dict=dict()
for pm in pdq:
    for pm_seasonal in seasonal_pdq:
        try:
            model=sm.tsa.statespace.SARIMAX(y_train['Sales'],
                                             order=pm,
                                             seasonal_order=pm_seasonal,
                                             enforce_stationarity=False,
                                             enforce_invertibility=False
                                             )

            model_fit=model.fit()
            print('SARIMA{X}{12} - AIC:{}'.format(pm,pm_seasonal,model_fit.aic
            ))
            metric_aic_dict.update({(pm,pm_seasonal):model_fit.bic})
        except:
            continue
```

SARIMA(0, 0, 0)X(0, 0, 0, 12)12 - AIC:810.3107683340173  
SARIMA(0, 0, 0)X(0, 0, 1, 12)12 - AIC:512.5318568681995  
SARIMA(0, 0, 0)X(0, 1, 0, 12)12 - AIC:465.20083573684155  
SARIMA(0, 0, 0)X(0, 1, 1, 12)12 - AIC:220.92269922440235  
SARIMA(0, 0, 0)X(1, 0, 0, 12)12 - AIC:486.90354702832656  
SARIMA(0, 0, 0)X(1, 0, 1, 12)12 - AIC:478.0942465724045  
SARIMA(0, 0, 0)X(1, 1, 0, 12)12 - AIC:250.99144098582636  
SARIMA(0, 0, 0)X(1, 1, 1, 12)12 - AIC:222.47824164007164  
SARIMA(0, 0, 1)X(0, 0, 0, 12)12 - AIC:760.2162374052343  
SARIMA(0, 0, 1)X(0, 0, 1, 12)12 - AIC:471.742008614918  
SARIMA(0, 0, 1)X(0, 1, 0, 12)12 - AIC:444.3094989817134  
SARIMA(0, 0, 1)X(0, 1, 1, 12)12 - AIC:203.57371601301926  
SARIMA(0, 0, 1)X(1, 0, 0, 12)12 - AIC:550.986052580633  
SARIMA(0, 0, 1)X(1, 0, 1, 12)12 - AIC:487.7711520891767  
SARIMA(0, 0, 1)X(1, 1, 0, 12)12 - AIC:251.80191286749474  
SARIMA(0, 0, 1)X(1, 1, 1, 12)12 - AIC:205.3034218803646  
SARIMA(0, 1, 0)X(0, 0, 0, 12)12 - AIC:728.2092083597857  
SARIMA(0, 1, 0)X(0, 0, 1, 12)12 - AIC:459.20277790229835  
SARIMA(0, 1, 0)X(0, 1, 0, 12)12 - AIC:450.79369617745346  
SARIMA(0, 1, 0)X(0, 1, 1, 12)12 - AIC:202.77719782465013  
SARIMA(0, 1, 0)X(1, 0, 0, 12)12 - AIC:468.58035819736926  
SARIMA(0, 1, 0)X(1, 0, 1, 12)12 - AIC:447.1273483957379  
SARIMA(0, 1, 0)X(1, 1, 0, 12)12 - AIC:233.31991014924887  
SARIMA(0, 1, 0)X(1, 1, 1, 12)12 - AIC:205.43611944734752  
SARIMA(0, 1, 1)X(0, 0, 0, 12)12 - AIC:704.1480775842399  
SARIMA(0, 1, 1)X(0, 0, 1, 12)12 - AIC:438.9991433259569  
SARIMA(0, 1, 1)X(0, 1, 0, 12)12 - AIC:430.16598457879934  
SARIMA(0, 1, 1)X(0, 1, 1, 12)12 - AIC:183.35240821973133  
SARIMA(0, 1, 1)X(1, 0, 0, 12)12 - AIC:465.1789965200763  
SARIMA(0, 1, 1)X(1, 0, 1, 12)12 - AIC:437.5033281512461  
SARIMA(0, 1, 1)X(1, 1, 0, 12)12 - AIC:227.09919145406244  
SARIMA(0, 1, 1)X(1, 1, 1, 12)12 - AIC:186.72632791183227  
SARIMA(1, 0, 0)X(0, 0, 0, 12)12 - AIC:750.1182176042614  
SARIMA(1, 0, 0)X(0, 0, 1, 12)12 - AIC:480.6501650967293  
SARIMA(1, 0, 0)X(0, 1, 0, 12)12 - AIC:463.48112766037684  
SARIMA(1, 0, 0)X(0, 1, 1, 12)12 - AIC:222.34642879389077  
SARIMA(1, 0, 0)X(1, 0, 0, 12)12 - AIC:479.2475854304206  
SARIMA(1, 0, 0)X(1, 0, 1, 12)12 - AIC:479.98268978067205  
SARIMA(1, 0, 0)X(1, 1, 0, 12)12 - AIC:224.8283373225528  
SARIMA(1, 0, 0)X(1, 1, 1, 12)12 - AIC:223.9439680897846  
SARIMA(1, 0, 1)X(0, 0, 0, 12)12 - AIC:724.6869847459585  
SARIMA(1, 0, 1)X(0, 0, 1, 12)12 - AIC:458.51987038484543  
SARIMA(1, 0, 1)X(0, 1, 0, 12)12 - AIC:446.1357737985943  
SARIMA(1, 0, 1)X(0, 1, 1, 12)12 - AIC:205.56596640901193  
SARIMA(1, 0, 1)X(1, 0, 0, 12)12 - AIC:475.5533631945233  
SARIMA(1, 0, 1)X(1, 0, 1, 12)12 - AIC:457.03110618710144  
SARIMA(1, 0, 1)X(1, 1, 0, 12)12 - AIC:226.58981409787648  
SARIMA(1, 0, 1)X(1, 1, 1, 12)12 - AIC:207.28234128817883  
SARIMA(1, 1, 0)X(0, 0, 0, 12)12 - AIC:727.5705796560088  
SARIMA(1, 1, 0)X(0, 0, 1, 12)12 - AIC:460.32096160056415  
SARIMA(1, 1, 0)X(0, 1, 0, 12)12 - AIC:451.7896977934569  
SARIMA(1, 1, 0)X(0, 1, 1, 12)12 - AIC:205.55468759202444  
SARIMA(1, 1, 0)X(1, 0, 0, 12)12 - AIC:459.212871336873  
SARIMA(1, 1, 0)X(1, 0, 1, 12)12 - AIC:459.51172759285146  
SARIMA(1, 1, 0)X(1, 1, 0, 12)12 - AIC:210.24225011225886  
SARIMA(1, 1, 0)X(1, 1, 1, 12)12 - AIC:207.4126038014479  
SARIMA(1, 1, 1)X(0, 0, 0, 12)12 - AIC:699.1410793411434

```
SARIMA(1, 1, 1)X(0, 0, 1, 12)12 - AIC:437.0892335544464
SARIMA(1, 1, 1)X(0, 1, 0, 12)12 - AIC:430.60551465582694
SARIMA(1, 1, 1)X(0, 1, 1, 12)12 - AIC:186.33341905894557
SARIMA(1, 1, 1)X(1, 0, 0, 12)12 - AIC:456.0424032256515
SARIMA(1, 1, 1)X(1, 0, 1, 12)12 - AIC:436.4921606265953
SARIMA(1, 1, 1)X(1, 1, 0, 12)12 - AIC:209.23315518105352
SARIMA(1, 1, 1)X(1, 1, 1, 12)12 - AIC:185.71420846266625
```

```
In [ ]: # The warnings means basically, the covariates are so big that the model has a
        # pretty hard time trying to fit coefficients for the linear regression and ever
        # ything else gets messed up.
        # https://stats.stackexchange.com/questions/340054/scaling-control-time-series-with-causalimpact
```

```
In [38]: {k: v for k,v in sorted(metric_aic_dict.items(),key=lambda x: x[1])}
```

```

Out[38]: {((0, 1, 1), (0, 1, 1, 12)): 183.94408195173997,
          ((1, 1, 1), (1, 1, 1, 12)): 186.70033134934735,
          ((1, 1, 1), (0, 1, 1, 12)): 187.12231736829045,
          ((0, 1, 1), (1, 1, 1, 12)): 187.51522622117716,
          ((0, 1, 0), (0, 1, 1, 12)): 203.38236801063823,
          ((0, 0, 1), (0, 1, 1, 12)): 204.4814712920014,
          ((0, 1, 0), (1, 1, 1, 12)): 206.34387472632966,
          ((1, 1, 0), (0, 1, 1, 12)): 206.46244287100657,
          ((0, 0, 1), (1, 1, 1, 12)): 206.51376225234077,
          ((1, 0, 1), (0, 1, 1, 12)): 206.7763067809881,
          ((1, 1, 0), (1, 1, 1, 12)): 208.62294417342406,
          ((1, 0, 1), (1, 1, 1, 12)): 208.79526675314906,
          ((1, 1, 1), (1, 1, 0, 12)): 210.4434955530297,
          ((1, 1, 0), (1, 1, 0, 12)): 211.150005391241,
          ((0, 0, 0), (0, 1, 1, 12)): 221.71848976999908,
          ((1, 0, 0), (0, 1, 1, 12)): 223.5401146122859,
          ((0, 0, 0), (1, 1, 1, 12)): 223.67192745846677,
          ((1, 0, 0), (1, 1, 1, 12)): 225.53554918097808,
          ((1, 0, 0), (1, 1, 0, 12)): 226.02202314094794,
          ((1, 0, 1), (1, 1, 0, 12)): 228.18139518906997,
          ((0, 1, 1), (1, 1, 0, 12)): 228.29287727245756,
          ((0, 1, 0), (1, 1, 0, 12)): 234.1157006948456,
          ((0, 0, 0), (1, 1, 0, 12)): 251.96125428540236,
          ((0, 0, 1), (1, 1, 0, 12)): 253.25663281685874,
          ((0, 1, 1), (0, 1, 0, 12)): 432.2550294542462,
          ((1, 1, 1), (0, 1, 0, 12)): 433.7390819689972,
          ((1, 1, 1), (0, 0, 1, 12)): 441.26732330534014,
          ((0, 1, 1), (1, 0, 1, 12)): 441.6814179021398,
          ((1, 1, 1), (1, 0, 1, 12)): 441.7147728152124,
          ((0, 1, 1), (0, 0, 1, 12)): 442.13271063912714,
          ((0, 0, 1), (0, 1, 0, 12)): 446.49158388843,
          ((1, 0, 1), (0, 1, 0, 12)): 449.40890115866927,
          ((0, 1, 0), (1, 0, 1, 12)): 450.40047575581286,
          ((0, 1, 0), (0, 1, 0, 12)): 451.8847386308118,
          ((1, 1, 0), (0, 1, 0, 12)): 453.97178270017355,
          ((1, 1, 1), (1, 0, 0, 12)): 460.4065730390848,
          ((0, 1, 0), (0, 0, 1, 12)): 461.384862809015,
          ((1, 1, 0), (1, 0, 0, 12)): 462.48599869694795,
          ((1, 0, 1), (1, 0, 1, 12)): 462.48631845389303,
          ((1, 0, 1), (0, 0, 1, 12)): 462.8840401982787,
          ((1, 1, 0), (0, 0, 1, 12)): 463.5940889606391,
          ((1, 1, 0), (1, 0, 1, 12)): 463.8758974062847,
          ((1, 0, 0), (0, 1, 0, 12)): 465.75211609223516,
          ((0, 0, 0), (0, 1, 0, 12)): 466.3363299527707,
          ((0, 1, 1), (1, 0, 0, 12)): 468.58547916786375,
          ((0, 1, 0), (1, 0, 0, 12)): 470.8513466292276,
          ((0, 0, 1), (0, 0, 1, 12)): 475.015135974993,
          ((1, 0, 1), (1, 0, 0, 12)): 480.0953400582399,
          ((0, 0, 0), (1, 0, 1, 12)): 481.50072922019194,
          ((1, 0, 0), (1, 0, 0, 12)): 482.65406807820807,
          ((1, 0, 0), (0, 0, 1, 12)): 484.0566477445168,
          ((1, 0, 0), (1, 0, 1, 12)): 484.5246666443886,
          ((0, 0, 0), (1, 0, 0, 12)): 489.25965468902245,
          ((0, 0, 1), (1, 0, 1, 12)): 492.13532190261,
          ((0, 0, 0), (0, 0, 1, 12)): 514.8028453000578,
          ((0, 0, 1), (1, 0, 0, 12)): 554.5202140716768,
          ((1, 1, 1), (0, 0, 0, 12)): 703.6306020255428,

```

```
((0, 1, 1), (0, 0, 0, 12)): 707.1410927071729,
((1, 0, 1), (0, 0, 0, 12)): 729.266066319807,
((0, 1, 0), (0, 0, 0, 12)): 729.7355688844019,
((1, 1, 0), (0, 0, 0, 12)): 730.6233007052411,
((1, 0, 0), (0, 0, 0, 12)): 753.2289137272403,
((0, 0, 1), (0, 0, 0, 12)): 763.2689584544667,
((0, 0, 0), (0, 0, 0, 12)): 811.8661163955068}
```

```
In [ ]: # We will compare models that better fit with pdq until 1, 2 and 3
```

```
In [43]: model=sm.tsa.statespace.SARIMAX(y_train['Sales'],order=(0,1,1),seasonal_order=
(0,1,1,12))
model_fit1=model.fit()
model_fit1.summary()
```

Out[43]: SARIMAX Results

<b>Dep. Variable:</b>	Sales	<b>No. Observations:</b>	36
<b>Model:</b>	SARIMAX(0, 1, 1)x(0, 1, 1, 12)	<b>Log Likelihood</b>	-232.841
<b>Date:</b>	Sun, 20 Dec 2020	<b>AIC</b>	471.681
<b>Time:</b>	18:19:06	<b>BIC</b>	475.088
<b>Sample:</b>	01-01-1997	<b>HQIC</b>	472.538
	- 12-01-1999		
<b>Covariance Type:</b>	opg		

	coef	std err	z	P> z	[0.025	0.975]
<b>ma.L1</b>	-0.4063	0.178	-2.286	0.022	-0.754	-0.058
<b>ma.S.L12</b>	-0.2952	0.200	-1.476	0.140	-0.687	0.097
<b>sigma2</b>	4.368e+07	1.33e-09	3.28e+16	0.000	4.37e+07	4.37e+07

<b>Ljung-Box (Q):</b>	9.74	<b>Jarque-Bera (JB):</b>	0.54
<b>Prob(Q):</b>	0.99	<b>Prob(JB):</b>	0.76
<b>Heteroskedasticity (H):</b>	1.92	<b>Skew:</b>	-0.37
<b>Prob(H) (two-sided):</b>	0.38	<b>Kurtosis:</b>	2.85

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

[2] Covariance matrix is singular or near-singular, with condition number 1.32e+32. Standard errors may be unstable.

```
In [39]: model=sm.tsa.statespace.SARIMAX(y_train['Sales'],order=(0,0,2),seasonal_order=(0,2,0,12))
model_fit2=model.fit()
model_fit2.summary()
```

Out[39]: SARIMAX Results

<b>Dep. Variable:</b>	Sales	<b>No. Observations:</b>	36
<b>Model:</b>	SARIMAX(0, 0, 2)x(0, 2, [], 12)	<b>Log Likelihood</b>	-124.712
<b>Date:</b>	Sun, 20 Dec 2020	<b>AIC</b>	255.425
<b>Time:</b>	18:18:40	<b>BIC</b>	256.879
<b>Sample:</b>	01-01-1997 - 12-01-1999	<b>HQIC</b>	254.886
<b>Covariance Type:</b>	opg		

	coef	std err	z	P> z	[0.025	0.975]
<b>ma.L1</b>	0.2829	0.233	1.216	0.224	-0.173	0.739
<b>ma.L2</b>	0.0461	0.225	0.205	0.838	-0.394	0.487
<b>sigma2</b>	6.203e+07	5.55e-10	1.12e+17	0.000	6.2e+07	6.2e+07

<b>Ljung-Box (Q):</b>	8.29	<b>Jarque-Bera (JB):</b>	1.58
<b>Prob(Q):</b>	0.69	<b>Prob(JB):</b>	0.45
<b>Heteroskedasticity (H):</b>	0.16	<b>Skew:</b>	0.86
<b>Prob(H) (two-sided):</b>	0.10	<b>Kurtosis:</b>	3.44

Warnings:

- [1] Covariance matrix calculated using the outer product of gradients (complex-step).
- [2] Covariance matrix is singular or near-singular, with condition number 1.39e+34. Standard errors may be unstable.



```
In [41]: model=sm.tsa.statespace.SARIMAX(y_train['Sales'],order=(0,3,3),seasonal_order=(0,2,0,12))
model_fit3=model.fit()
model_fit3.summary()
```

Out[41]:

SARIMAX Results

<b>Dep. Variable:</b>	Sales	<b>No. Observations:</b>	36
<b>Model:</b>	SARIMAX(0, 3, 3)x(0, 2, [], 12)	<b>Log Likelihood</b>	-93.005
<b>Date:</b>	Sun, 20 Dec 2020	<b>AIC</b>	194.010
<b>Time:</b>	18:18:50	<b>BIC</b>	194.799
<b>Sample:</b>	01-01-1997 - 12-01-1999	<b>HQIC</b>	192.308
<b>Covariance Type:</b>	opg		

	coef	std err	z	P> z	[0.025	0.975]
<b>ma.L1</b>	-2.1833	0.869	-2.511	0.012	-3.887	-0.479
<b>ma.L2</b>	1.4463	1.323	1.093	0.274	-1.147	4.040
<b>ma.L3</b>	-0.2382	0.646	-0.369	0.712	-1.503	1.027
<b>sigma2</b>	5.165e+07	4.87e-08	1.06e+15	0.000	5.16e+07	5.16e+07

<b>Ljung-Box (Q):</b>	3.70	<b>Jarque-Bera (JB):</b>	0.49
<b>Prob(Q):</b>	0.88	<b>Prob(JB):</b>	0.78
<b>Heteroskedasticity (H):</b>	0.21	<b>Skew:</b>	-0.31
<b>Prob(H) (two-sided):</b>	0.23	<b>Kurtosis:</b>	2.04

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).


[2] Covariance matrix is singular or near-singular, with condition number 1.15e+31. Standard errors may be unstable.

```
In [ ]: #From the three Sarima models, we see that the one that has the best performance is the third one.
#This is logical because we have differentiated 3 lags
```

## 7. Comparing results

```
In [44]: from datetime import datetime
df['forecast1']=model_fit1.predict(start=datetime(2000,1,1), end=datetime(2000,12,1),dynamic=True)
df['forecast2']=model_fit2.predict(start=datetime(2000,1,1), end=datetime(2000,12,1),dynamic=True)
df['forecast3']=model_fit3.predict(start=datetime(2000,1,1), end=datetime(2000,12,1),dynamic=True)
df[['Sales','forecast1','forecast2','forecast3']].plot(figsize=(12,8))
```

Out[44]: <matplotlib.axes.\_subplots.AxesSubplot at 0x224afa2ada0>



```
In [45]: actual=y_test['Sales']['2000-01-01':]
forecast1=model_fit1.get_prediction(start=pd.to_datetime('2000-01-01'),end=pd.to_datetime('2000-12-01'),dynamic=False)
predictions1=forecast1.predicted_mean
rmse1=np.sqrt((predictions1 - actual)**2).mean()

forecast2=model_fit2.get_prediction(start=pd.to_datetime('2000-01-01'),end=pd.to_datetime('2000-12-01'),dynamic=False)
predictions2=forecast2.predicted_mean
rmse2=np.sqrt((predictions2 - actual)**2).mean()

forecast3=model_fit3.get_prediction(start=pd.to_datetime('2000-01-01'),end=pd.to_datetime('2000-12-01'),dynamic=False)
predictions3=forecast3.predicted_mean
rmse3=np.sqrt((predictions3 - actual)**2).mean()
print('Rmse1: {}'.format(rmse1))
print('Rmse2: {}'.format(rmse2))
print('Rmse3: {}'.format(rmse3))
```

```
Rmse1: 6495.958370396597
Rmse2: 6208.075242555135
Rmse3: 14346.355962425181
```

```
In [ ]: #We can see that even aic is better for sarima model3, the one that fits best
        with the test part is the second model
        # It is maybe because the testing year doesn't have the same trend as the previous year
```

## 8. Forecasting with new model


```
In [46]: from pandas.tseries.offsets import DateOffset
future_dates=[df.index[-1]+ DateOffset(months=x)for x in range(0,12)]
```

```
In [47]: future_datest_df=pd.DataFrame(index=future_dates[1:],columns=df.columns)
```

```
In [48]: future_df=pd.concat([df,future_datest_df])
```

```
In [49]: future_df['forecast1'] = model_fit1.predict(start = datetime(2000,1,1), end =
datetime(2003,1,1), dynamic= True)
future_df['forecast2'] = model_fit2.predict(start = datetime(2000,1,1), end =
datetime(2003,1,1), dynamic= True)
future_df['forecast3'] = model_fit3.predict(start = datetime(2000,1,1), end =
datetime(2003,1,1), dynamic= True)
future_df[['Sales', 'forecast1', 'forecast2', 'forecast3']].plot(figsize=(12, 8
))
```

Out[49]: <matplotlib.axes.\_subplots.AxesSubplot at 0x224afb140b8>



```
In [ ]: #this graphic finally gives us a view in which we can see that model3 is not good at all. The ideal model would be a combination between the first and the second model
#the third model is bad, maybe because it is detecting a descending trend
```